Customer Quality Index (CQI) — Full Product Blueprint

Purpose

The Customer Quality Index (CQI) is a zero-auth, Shopify-only diagnostic micro-SaaS for D2C founders and CEOs. Its job is to use public data plus a short calibration intake to produce an institutional-looking audit that feels worth at least $5,000, reveals hidden profit leaks, and converts qualified operators into agency opportunities.[web:34] [web:37] [web:31]

The product is not positioned as generic website auditing software. It is positioned as a growth-side forensic system that helps founders sanity-check their ad spend, detect structural waste, and understand how public storefront signals may be suppressing both near-term profitability and long-term equity value.[web:1] [web:3] [web:8]

Product definition

Core promise

CQI tells a founder: your customer acquisition engine may be attracting the wrong buyer, training ad networks on the wrong signals, wasting spend through stock and tracking gaps, and lowering the quality of your asset profile.[web:1] [web:8]

The platform does this without logins, private API tokens, Shopify admin access, or ad account permissions. It combines Shopify storefront scraping, Meta Ad Library observation, public tech-stack detection, public review extraction, schema visibility analysis, and Google AI Mode style grounding through Apify actors only.[web:6] [web:28] [web:31] [web:34] [web:37] [web:42]

Commercial role

CQI is not the cure. It is the evidence engine that makes the agency cure feel necessary, justified, and already partially designed. That means every report output has two jobs:

- Prove something is leaking.
- Show the repaired future state as specific and financially meaningful.

User definition

Primary user:

- Founder or CEO of a Shopify D2C brand.

Target brand profile:

- Mostly $500k to $15M ARR brands.
- Cross-category, but always normalized into a category-specific benchmark model.
- Shopify only.

Primary emotional job-to-be-done:

- “Find my hidden leaks and sanity-check my ad spend.”

Secondary emotional subtext:

- “I suspect my agency or internal media team is overstating reality.”
- “I want an independent, clinical read on my store without giving access.”
- “I want something I can forward internally and use as leverage.”

Product principles

Principle 1: Zero-friction authority

The product must feel advanced before asking for anything. No login, no email gate, no account creation, no payment wall. This lowers friction and raises curiosity.

Principle 2: Institutional tone, growth relevance

The tone is not loud or agency-cheap. It uses sober, clinical language and financial framing, but the content is practical and growth-oriented rather than banker-only.

Principle 3: Evidence first, alarm second

The founder should first see objective operational evidence, then financial consequence. This avoids instant defensiveness.

Principle 4: Every weakness must imply a paid implementation path

Every red zone in the report must naturally map to one agency protocol. Never leave a major defect hanging without a named future-state resolution.

Principle 5: Fallbacks must preserve dignity

If public data is thin, the software should never look broken. It should shift into clearly labeled simulation, opacity, or inferred-baseline modes rather than fail awkwardly.

Scope boundaries

Included

- Shopify brands only.
- Public storefront data.
- Public review data.
- Public schema and crawlability.
- Public Meta Ad Library visibility when available.
- Public tech footprint detection.
- Search-grounded market context through Google AI Mode actors on Apify.[web:31] [web:36]

Excluded

- Non-Shopify brands.
- Private ad account data.
- Private Shopify admin data.
- Email/SMS backend performance data.
- Agency-style multi-user workspace features.
- User accounts and persistent cloud profiles.

User journey

Stage 1: Landing and framing

The user lands on a highly polished, high-contrast interface that looks like an institutional audit terminal, not a landing page template. The page must immediately communicate three beliefs:

- This is Shopify-native.
- This is independent and zero-access.
- This is about capital efficiency, not vanity metrics.

Above-the-fold message architecture

Primary headline:

- Audit the quality of your customer acquisition engine.

Primary subhead:

- Detect hidden profit leakage, signal distortion, inventory friction, and AI visibility gaps using only your public storefront footprint.

Trust frame components:

- “No login.”
- “No Shopify access.”
- “No ad account connection.”
- “Shopify brands only.”
- “Baseline advisory value: $5,000 | Sponsored allocation.”

Psychology at this stage:

- Reduce fear of giving access.
- Increase curiosity.
- Signal prestige.
- Pre-anchor the report as expensive.

Stage 2: Intake console

The intake panel is the only required input moment before report generation.

Inputs

1. Target Shopify domain.
2. Total combined monthly paid performance budget.
3. Average order value.
4. Estimated storewide return rate.
5. Promotional code checkout frequency.
6. Primary vertical dropdown.
7. Optional exact ad spend override.
8. Optional competitor domain for later benchmarking input, either before run or inside report.

Input behavior

- Domain input strips protocol and subpaths.
- Shopify detection runs immediately in background.
- If non-Shopify is detected, user sees a graceful unsupported message.
- Spend uses a fast slider plus optional exact input toggle.
- All financial inputs are normalized into USD-like internal values.
- Inputs are bracket-friendly, but exact numbers override bracket averages.

Input copy style

Every input label should sound operator-grade, not marketing-grade. Example:

- “Total combined monthly budget deployed across paid performance networks targeting digital purchase pixels.”

Intake psychology

The intake is not just data collection. It primes the founder to think in terms of:

- paid systems,
- leakage,
- return drag,
- discount dependency,
- category benchmarking.

That means every field should quietly reinforce the report’s eventual logic.

Processing architecture

High-level orchestration

The orchestration starts as soon as the user clicks the compute button.

Input:

- Shopify domain.
- 4–8 calibration values.

Processing branches:

- Quantitative extraction layer.
- Qualitative cognitive layer.
- Normalization layer.
- Deterministic scoring and financial model layer.
- Report assembly layer.
- Cache and export layer.

Live processing screen

During the 60–90 second run, the UI shows a real-time execution log. This is both functional and psychological.

Purpose of the execution log:

- Demonstrates that the audit is doing real work.

- Creates perceived sophistication.
- Justifies the eventual authority of the report.
- Prevents the wait from feeling broken.

Sample execution messages:

- Initializing anonymous public footprint routing.
- Validating Shopify storefront architecture.
- Pulling products and variant availability maps.
- Mapping Meta destination strings and creative velocity.
- Reading schema graphs for AI crawler visibility.
- Extracting customer review friction themes.
- Computing category-adjusted risk weights.
- Rendering capital efficiency matrix.

Data pipeline

Input layer

User-supplied data

Stored immediately in a session object:

```
{
  "domain": "brand.com",
  "spend_input": 50000,
  "spend_exact": 430000,
  "aov": 110,
  "return_rate": 0.20,
  "promo_frequency": 0.45,
  "vertical": "Apparel"
}
```

Public extraction layer

Source 1: Shopify product layer

Primary method:

- Shopify product scraper actor or direct products feed extraction through an Apify Shopify actor.[web:34] [web:35] [web:37]

Expected outputs:

- Product handles.

- Titles.
- Collections.
- Images.
- Variants.
- Variant availability.
- Product URLs.
- Collection paths.

Purpose:

- Inventory conflict analysis.
- Product density estimation.
- Best-seller approximation support.
- Evidence for ad destination mismatch.

Fallback:

- PLP crawl through Shopify collection pages if exact product/variant visibility is constrained.

Source 2: Meta Ad Library layer

Primary method:

- Meta/Facebook Ad Library scraper actor on Apify.[web:6] [web:28]

Expected outputs:

- Active ad count.
- Creative count.
- Start dates or creative age when available.
- Destination URLs.
- Publisher/account metadata when available.

Purpose:

- Creative volume and testing velocity.
- Destination path mapping.
- Signal opacity or dark-loop state detection.
- Evidence for active traffic being sent to weak product or collection pages.

Fallback state:

- “Signal Opaque / Dark Loop Status.”

Source 3: Tech stack footprint layer

Primary method:

- Cheapest stable actor that detects Shopify storefront technologies and script signatures.

Additional custom pass:

- HTML/script scan for window variables, script URLs, and object signatures indicating tools like Triple Whale, Elevar, Northbeam, Klaviyo, review apps, schema tools, and tracking containers.

Purpose:

- Tracking maturity analysis.
- SaaS redundancy analysis.
- Pixel/CAPI architecture clues.
- Review widget detection.
- AI visibility clues.

Source 4: Review and sentiment layer

Primary method:

- Google AI Mode / Gemini-style extraction via Apify actor, pointed first at onsite review widgets and product pages.[web:31] [web:36]

Review hierarchy:

1. On-site review widget content.
2. External review platforms like Trustpilot or Google business review surfaces when relevant.
3. Public discussion nodes like Reddit if needed.

Expected outputs:

- Review text excerpts.
- Dominant complaint clusters.
- Sentiment score.
- Product-level friction themes.
- Fulfillment, sizing, quality, or expectation mismatch markers.

Purpose:

- Public customer quality proxy.
- Friction pattern extraction.
- Copy-message mismatch analysis.
- Valuation risk narrative support.

Source 5: Schema and AI visibility layer

Primary method:

- Crawl homepage, PDPs, collection pages, and robots/schema surfaces.
- Use Apify Google AI Mode actor and direct HTML parsing to assess crawler-visible text, JSON-LD, review readability, and semantic graph presence.[web:31] [web:42]

Expected outputs:

- JSON-LD presence.
- Product schema quality.
- Review text crawlability.
- robots accessibility clues.
- AI readability index.

Purpose:

- Module 4.
- High-performing brand wedge.
- Future-proofing narrative.

Source 6: Macro and category baseline layer

Primary method:

- Google AI Mode / Gemini-related actors on Apify for live category norms, category-specific cost logic, market narratives, and valuation multiple context.[web:31] [web:36]

Expected outputs:

- Category CPC/CVR benchmark support.
- Margin expectation clues.
- Bear/base/bull multiple references.
- Sector-specific macro narrative.

Purpose:

- Baseline constants.
- Financial model support.
- Report narrative credibility.

Evidence sufficiency logic

CQI uses a 3-of-5 evidence threshold.

Core evidence vectors:

- Meta Ad Library.
- On-site reviews.
- Tech stack scripts.
- JSON-LD/schema graphs.
- Accessible PLP structures.

If at least 3 are available:

- Run full evidence-weighted audit.

If fewer than 3 are available:

- Switch to Simulation & Schema Integrity Audit.
- Label the report clearly as simulated from partial public signals and category baselines.

This preserves authority by turning missing data into a professional state, not an error.

Normalization layer

Normalization role

All raw data from actors and scrapers arrives inconsistent. The normalization layer converts this into a single structured JSON object. In the current design, raw public signals are transformed into clean semantic payloads before deterministic calculation.

Canonical normalized object

```
{
  "target_store": { "domain": "brand.com", "platform": "shopify" },
  "sliders": { "ad_spend": 50000, "aov": 110, "return_rate": 0.20, "promo_frequency
  "brand_profile": { "persona_avatar": "String", "visual_vibe": "String", "sentime
  "quantitative_metrics": { "active_meta_ads": 14, "variant_stockout_ratio": 0.35, "
  "tech_stack": { "detected_apps": ["Triple Whale", "Elevar", "Klaviyo"], "spread_l
  "ai_readability": { "schema_valid": false, "crawlable_reviews": false, "index_sco
  "macro_multiples": { "category_vertical": "Apparel", "target_multiple": 3.8, "re
  "valuation_risk_matrix": { "ideal_value": 4850000, "markdown_penalty": 42, "true_v
}
```

Extended internal object requirements

The production object should also include:

- `run_id`
- `evidence_confidence_score`
- `data_mode` (`full`, `partial`, `simulated`)
- `review_sources`
- `meta_signal_state`
- `inventory_visibility_state`
- `vertical_selected`
- `vertical_refined`
- `module_flags`
- `cta_mappings`
- `pdf_render_ready`

Scoring engine

CQI vectors

The master score uses six core vectors:

1. Technical Tracking Integrity.
2. Variant Stockout Exposure.
3. Promotional Coupon Dependency.
4. AI Crawler Visibility.
5. Creative Volume and Testing Velocity.
6. Public Customer Review Sentiment.

Weighting logic

The vectors remain globally consistent, but the weights are vertical-adjusted through a Supabase config table. This keeps the schema simple while aligning the model to real category structure.

Example:

- Apparel may tolerate higher returns.
- Beauty may weight repeatability and sentiment more heavily.
- Digital products may set return-rate effects to zero and reassign weight to tracking and AI visibility.

Score labels

Global label bands remain fixed:

- 85–100: Asset Compounder.
- 70–84: Exposed Operator.
- 55–69: Fragile Growth.
- 40–54: Ad Treadmill.
- Below 40: Margin Leak Engine.

Benchmark display

The interface shows both:

- the brand's score,
- the vertical baseline.

This converts the result from an isolated opinion into a competitive context.

Financial model

Purpose of the model

The financial model does not claim forensic accounting precision. It translates public technical and behavioral signals into a founder-readable estimate of commercial waste and valuation sensitivity.

Core equations

Annual revenue inference:

$$R_{annual} = \left(\frac{S_{monthly}}{CPC_{category}} \right) \times CR_{category} \times AOV \times 12$$

Ideal valuation:

$$V_{ideal} = R_{annual} \times M_{category}$$

Compounded risk discount:

$$D_{risk} = (\max(0, TargetMargin - TrueMargin) \times 100 \times 3.0) + (O_{variant} \times 100 \times 0.5$$

Equity loss:

$$L_{equity} = V_{ideal} \times \left(\frac{D_{risk}}{100} \right)$$

Guardrails

- Revenue floor for institutional mode: $1M baseline asset scale.
- Risk discount cap: 85%.
- Bear/Base/Bull shown as a band, with Base highlighted.
- Vertical-specific multiple table, updated through config.

Micro-brand handling

If inferred scale is below $1M:

- Remove traditional PE-style valuation panel.
- Replace with “Acquisition Ready Revenue Graduation Horizon.”
- Compare future $1M scale under dirty vs clean infrastructure conditions.

This turns small size into a roadmap instead of a credibility problem.

Report structure

Opening layer

The report opens on an executive summary, not a score shock.

Executive summary components

1. Headline verdict.
2. 3–4 key bullet insights.
3. Data certainty score.
4. Vertical benchmark note.
5. Sponsored allocation / advisory value anchor.

Psychology

This creates trust before anxiety. The founder first sees that the system has found concrete truths, then sees the commercial consequence.

Section order

1. Executive Summary.
2. CQI Master Grade.
3. Risk Fragility Alert.
4. Five deep-dive modules.
5. Competitive Shadow Benchmark.
6. Institutional valuation band.

7. Resolution Protocol CTAs.
8. Methodology drawer.
9. PDF export.

Module design

Common module frame

Every module follows the same cognitive structure:

1. Name of the leak.
2. What public evidence was observed.
3. Current bleeding state.
4. Future optimized state.
5. Estimated financial or strategic consequence.
6. Single flagship CTA.

This consistency trains the founder to think in terms of diagnosed problem followed by engineered resolution.

Module 1: Ad Signal Contamination Matrix

Objective

Show that ad networks may be optimizing toward low-quality buyers because the business is feeding the wrong value logic into acquisition systems.

Inputs

- User spend.
- User promo frequency.
- Category norms.
- Tech stack clues.
- Review quality and brand positioning cues.
- Ad library visibility.

Logic

Estimate distortion between reported checkout value and likely contribution-quality buyer value.

Output examples

- Passed Pixel Value.
- Estimated True Unit Profit.
- Signal Distortion Gap.
- High-discount segment contamination note.

Psychology

This is often the strongest wedge because it validates the founder's suspicion that ad performance reporting is misleading.

CTA

- Review Profit Visibility Engine Protocol.

Module 2: Best-Seller Stockout Conflict Ledger

Objective

Show that paid traffic may be routed toward degraded or unavailable product experiences.

Inputs

- Product/variant availability.
- Ad destinations.
- Collection page signals.
- Review-weighted flagship product clues.

Logic

Map active ad destinations to stock-poor or variant-poor pages and estimate wasted paid attention.

Outputs

- Top 3 conflict URLs.
- Variant deficit share.
- Estimated wasted ad burn.
- Inventory visibility state.

CTA

- Access Inventory Synchronization Architecture.

Module 3: Tech Stack Redundancy Paradox

Objective

Show that the brand pays for serious software but leaves it operationally underconfigured.

Inputs

- Script signatures.
- Detected apps.
- Tracking stack count.
- Gaps surfaced in Modules 1 and 2.

Logic

The more advanced the detected stack and the worse the observed leaks, the stronger the redundancy paradox.

Outputs

- Detected premium tools.
- Estimated annual software waste.
- Overlap/conflict note.
- Consolidation opportunity narrative.

CTA

- View Structural Stack Optimization Roadmap.

Module 4: Conversational Engine Invisibility Index

Objective

Measure whether the brand is legible to AI-driven discovery and recommendation systems.

Inputs

- JSON-LD.
- Review crawlability.
- Product schema quality.
- Plain-text evidence visibility.

- robots clues.
- AI search visibility actor support if used.[web:42]

Logic

A clean, crawlable, semantically rich storefront receives a higher AI readability score.

Outputs

- AI readability score.
- Crawlable review state.
- Schema validity flag.
- Discovery moat opportunity.

Psychology

This is the most important wedge for already strong brands. It creates a future-risk conversation rather than a rescue conversation.

CTA

- Access AI Discovery Layer Protocol.

Module 5: Terminal Enterprise Valuation Markdown

Objective

Aggregate the discovered leaks into a top-level asset-quality framing.

Inputs

- Output from all modules.
- Category multiples.
- Review friction themes.
- AI visibility penalties.
- Margin/risk penalties.

Logic

Translate operational disorder into implied market discount, while showing the clean-state upside.

Outputs

- Bear/Base/Bull valuation band.
- Estimated markdown percentage.
- Equity value lost.
- Clean-state upside.

CTA

- Review Capital Efficiency Blueprint.

Competitive Shadow Benchmark

Purpose

Make the product sharable, comparative, and harder to dismiss.

Flow

The founder enters one competitor domain. The same public evidence framework runs against that competitor.

Output structure

- Your brand vs competitor.
- Creative testing velocity.
- Review response maturity.
- AI visibility score.
- Data certainty score.
- Market disadvantage notes.

Thin-data handling

If competitor visibility is weak, present it as a hardened or obfuscated footprint and downgrade certainty, not authority.

Psychological system design

Core emotional arc

The CQI report should make the founder move through five emotional states in this order:

1. Curiosity.
2. Recognition.
3. Controlled concern.

4. Strategic imagination.
5. Action readiness.

How each screen contributes

Landing page

- Curiosity through exclusivity and zero-friction access.

Intake

- Commitment through self-disclosure and business framing.

Processing screen

- Confidence through visible machinery.

Executive summary

- Recognition through specific truths.

Modules

- Controlled concern through quantified leakage.

Future-state language

- Strategic imagination through “dream state” system repair.

CTAs

- Action readiness through single-path, high-status implementation options.

Language rules

Use:

- capital inefficiency,
- signal contamination,
- asset degradation,
- data architecture,
- structural leakage,
- market share leakage,
- visibility gap,
- contribution drag.

Avoid:

- panic phrasing,
- founder blame,
- “your business is broken,”
- “book a call now,”
- insolvency or layoffs.

UX map

Screen 1: Landing

Contains:

- Headline.
- Subhead.
- Sponsored advisory anchor.
- Trust bars.
- Intake panel.
- Shopify-only note.

Screen 2: Processing

Contains:

- Animated execution log.
- Estimated run time.
- Data source badges.
- Progress state labels.

Screen 3: Executive summary

Contains:

- Summary verdict.
- 3–4 bullets.
- Data certainty score.
- Vertical baseline reference.
- PDF export button.

Screen 4: CQI dashboard

Contains:

- CQI master grade.
- Risk alert card.
- Bear/Base/Bull strip.
- Module navigation.

Screen 5: Deep-dive modules

Each contains:

- current state,
- evidence,
- math graphic,
- dream state,
- CTA.

Screen 6: Competitive benchmark

Contains:

- rival compare table,
- certainty note,
- highlighted deficits.

Screen 7: Methodology drawer

Contains:

- formulas,
- variable definitions,
- evidence confidence explanation,
- simulation/fallback notices where applicable.

Data storage model

Persistence logic

One brand = one latest snapshot record for fast reuse.

Core behavior:

- If same brand is requested again within 24 hours, serve cache.

- Keep `run_id` for potential historical support.
- Save entire report payload including competitor benchmark inside the brand record.

Suggested schema

```
CREATE TABLE public.cqi_diagnostic_audits (
    shopify_url TEXT PRIMARY KEY,
    created_at TIMESTAMP WITH TIME ZONE DEFAULT TIMEZONE('utc'::text, NOW()),
    run_id TEXT,
    user_spend_tier NUMERIC NOT NULL,
    user_aov NUMERIC NOT NULL,
    user_return_rate NUMERIC NOT NULL,
    user_discount_frequency NUMERIC NOT NULL,
    vertical_selected TEXT,
    vertical_refined TEXT,
    cqi_master_grade INT4 NOT NULL,
    inferred_annual_revenue NUMERIC(15, 2),
    potential_enterprise_valuation NUMERIC(15, 2),
    calculated_risk_discount_percent NUMERIC(5, 2),
    terminal_equity_loss NUMERIC(15, 2),
    evidence_confidence_score NUMERIC(5, 2),
    data_mode TEXT,
    signal_contamination_metadata JSONB DEFAULT '{}'::jsonb,
    inventory_conflict_ledger JSONB DEFAULT '{}'::jsonb,
    tech_stack_footprint JSONB DEFAULT '{}'::jsonb,
    ai_visibility_metrics JSONB DEFAULT '{}'::jsonb,
    review_sentiment_payload JSONB DEFAULT '{}'::jsonb,
    competitor_benchmark_payload JSONB DEFAULT '{}'::jsonb,
    client_ip_address_hash TEXT NOT NULL,
    cache_expiration_timestamp TIMESTAMP WITH TIME ZONE NOT NULL
);
```

Actor selection strategy

Rule for choosing actors

Choose the cheapest actor that satisfies all of the following:

- Stable output shape.
- Enough run history or marketplace maturity to trust basic uptime.
- Inputs can be capped.
- Outputs contain only the fields needed for CQI.
- No unnecessary browser-heavy logic if a feed-based approach exists.

Priority by layer

Shopify products

Prefer a feed or API-style Shopify actor over a browser crawler because Shopify product data is often structurally accessible and cheaper to pull.[web:34] [web:35] [web:37]

Ad library

Choose a Meta Ad Library actor with destination URL support and ad-count visibility. Cap max ads fetched per run to keep spend predictable.[web:6] [web:28]

Google AI Mode

Use a Google AI Mode/Gemini-style actor only for the tasks where reasoning and web grounding matter:

- review synthesis,
- category refinement,
- macro multiple collection,
- PR/news grounding,
- copy tone assessment.[web:31] [web:36]

AI visibility

If the dedicated AI visibility actor is cheap enough, use it only for Module 4 enrichment rather than as a core dependency.[web:42]

Build order

Phase 1: Thin-slice prototype

Build only:

- Shopify validation.
- Intake.
- Product extraction.
- Ad extraction.
- Tech stack extraction.
- One normalized payload.
- CQI master grade.
- Module 1.
- Executive summary.

Reason:

- This creates an end-to-end proof of value before building every advanced narrative branch.

Phase 2: Full audit core

Add:

- Module 2.
- Module 3.
- Module 4.
- Module 5.
- Bear/Base/Bull.
- Vertical config system.
- Simulation mode.

Phase 3: Conversion and defensibility

Add:

- Competitive benchmark.
- PDF export.
- Resolution Protocol page.
- Cookie/session calibration save.
- Evidence confidence score.
- Better fallback messaging.

Failure-state design

Never show raw failure

Every failure becomes a productized state:

- Meta unavailable -> Signal Opaque / Dark Loop.
- Inventory hidden -> Inventory Coupling Override Active.
- Low evidence -> Cognitive Simulation Deployed.
- Thin competitor data -> Defensible Moat Asset / Low certainty.

Why this matters

A founder should never think “this tool broke.” They should think “this store is hard to read publicly, and that itself is diagnostically meaningful.”

PDF logic

Export behavior

PDF is generated only on explicit click.

Purpose:

- Save compute budget.
- Use export as a value moment.
- Encourage sharing internally.

PDF composition

The PDF should be deck-like and around 20 pages, containing:

- cover page,
- executive summary,
- scorecard,
- module pages,
- benchmark page,
- valuation page,
- methodology appendix.

Resolution Protocol bridge

Role

The Resolution Protocol page is the bridge from diagnosis to agency conversion.

Content blocks

- Named leak.
- What is happening technically.
- What the repaired system looks like.
- Typical implementation timeline.
- Why the solution is architecture, not media buying.
- High-status CTA to engage further.

Pricing posture

Do not show exact retainer pricing. Maintain premium mystery and force discussion into system design, not cost shopping.

Metrics for success

Product metrics

- Audit start rate.
- Audit completion rate.
- PDF export rate.
- Module CTA click-through rate.
- Competitor benchmark usage rate.
- Return visit rate via cached domain reload.

Agency metrics

- Resolution Protocol click-through.
- Qualified meeting rate from report viewers.
- Close rate by CQI band.
- Average deal size by module trigger pattern.

Final product definition

CQI is a Shopify-only, zero-auth, public-data D2C diagnostic engine that turns storefront, ad-library, review, schema, and tech-footprint signals into a category-adjusted founder audit. It is engineered to feel like a sponsored $5,000 advisory artifact, to expose capital inefficiency without private access, and to convert evidence-backed urgency into agency implementation demand through a controlled, high-status resolution path.[web:1] [web:3] [web:6] [web:28] [web:31] [web:34] [web:37] [web:42]